

Nap

User Manual

Michael Dungan

Nap Manual



Figure 1: Nap interface

What is Nap?

Nap is a character reverb built on the **Extended Dark Velvet Noise (EDVN)** engine. Its defining feature is that you draw three curves — **Decay**, **Width**, and **Tone** — over a shared tail-position axis (0 → 100 %) that sculpt exactly how the reverb tail's loudness, stereo image, and tonal colour each evolve along their length.

That means you can draw a reverse swell, a plateau that collapses, a tail that starts mono and fans into full stereo at the halfway point, or one that starts bright and dims as it decays — all with the same three curve editors. Drawing the tail's loudness, stereo width, and tonal colour as three independent curves over its own timeline is an unusually direct degree of per-section

control — most reverbs expose these as fixed knobs or per-band settings rather than shapes you draw across the tail.

The DSP is time-domain sparse convolution (no FFT) by default, so Nap reports **zero latency** to the host. An optional **Efficient** engine switches to uniformly-partitioned FFT convolution (UPOLS), reducing CPU cost at large/dense settings at the cost of a short algorithmic latency (~512 samples ≈ 10.7 ms at 48 kHz). Pre-Delay shifts only the wet path; the dry signal stays phase-aligned in both engine modes.

Inspired by the Aalto Acoustics Lab papers on EDVN (Fagerström et al., JAES 2024) and BDVN (DAFx 2024).

Installation

Build from source (requires stable Rust):

```
cargo nih-plugin bundle nap --release
```

The bundler outputs to target/bundled/. Copy either the .vst3 or .clap file to your plugin directory:

- **Linux:** ~/.vst3/ or ~/.clap/
- **macOS:** ~/Library/Audio/Plug-Ins/VST3/ or ~/Library/Audio/Plug-Ins/CLAP/
- **Windows:** C:\Program Files\Common Files\VST3\ or C:\Program Files\Common Files\CLAP\

Quick Start

1. Insert Nap on a track and play audio through it.
2. Set **Mix** to around 30 % to blend in the tail.
3. The default **Decay** curve falls from full to silence — the tail decays naturally. Try sketching a plateau or a reverse swell.
4. Raise the **Width** dial to spread the tail stereo image; the **Width curve** defaults to a narrow → wide ramp so the tail starts centred and blooms wide, like a real room. Reshape it to control how wide each section of the tail is.
5. The **Tone curve** defaults to bright-at-attack, darkening over time (air absorption). Flatten it or invert it to taste.
6. Adjust **Size** (tail length) and **Density** (pulse count) to change the reverb character. The live tail visualization in the Decay pane updates as you edit.
7. Drag the **Seed** dial to re-roll the random velvet pattern while keeping the curve shapes.

Controls

Curve Editors (top three panes)

The three stacked MSEG editors share a **tail-position x-axis** running from 0 % (attack of the tail) to 100 % (end of the tail). All three drive the velvet sequence generation together; any edit triggers an immediate re-generation.

Reading the curves. Every curve runs left-to-right along the same axis — the *left edge is the start of the tail* (just after the dry hit) and the *right edge is the end of the tail* (where it fades out). The y-axis is “more of that quality”: higher = louder (Decay), wider (Width), or brighter (Tone). The factory defaults are deliberately shaped to read as one **natural room**: the tail **decays** to silence, **widens** from a centred onset into a diffuse stereo bloom, and **darkens** as high frequencies are absorbed. That gives you a familiar starting point — drag any curve from there to sculpt the part of the tail you want to change.

Curve	Controls	Default shape (start → end)
Decay	Tail loudness $g(m)$	Full → silence (exponential fall)
Width	Stereo spread $J(m)$	Narrow → wide (0.15 → 0.85 ramp)
Tone	Brightness (filter routing)	Bright → dark (0.85 → 0.25)

Decay Controls the **energy gain envelope** of the tail. The y-axis maps directly to pulse gain $g(m)$ at each point along the tail.

- **1.0 (top)** → full-amplitude pulses at that tail position.
- **0.0 (bottom)** → silent pulses; the tail is completely dead at that section.
- Draw an **exponential fall** (default) for a natural-sounding room.
- Draw a **gate** (full, then sudden silence) for a gated reverb.
- Draw a **reverse curve** (0 → 1) for a swell-style tail.
- Draw a **plateau** or **multi-bump** shape for rhythmic reverb textures.

The live pulse-field visualization in this pane shows the actual generated sequence: stick height = pulse gain, colour = tone/filter index, horizontal L/R split = width jitter.

Width Controls the **stereo spread** of the tail, section by section. The y-axis maps to the per-pulse maximum jitter $J(m)$ applied to the right-channel pulse locations.

- **0.0 (bottom)** → L and R channels receive identical pulse locations → the tail collapses to mono at that point.
- **1.0 (top)** → maximum jitter (set by the **Width** dial) → widest spread at that point.
- Draw a **flat 0** for a fully mono tail; draw a **flat 1** for uniform stereo spread.
- Draw a **ramp** (the default: 0.15 → 0.85) to have the image open up progressively along the tail — narrow, centred early reflections blooming into a wide, diffuse late tail, as in a real room.
- At Width dial = 0 ms, the curve has no effect (mono throughout).

The width mechanism is grounded in EDVN/BDVN theory: interaural coherence equals the DTFT of the pulse-location jitter distribution. Animating jitter width over the tail animates the stereo image at zero additional compute cost.

Tone Controls the **spectral colour** of each section of the tail. The y-axis selects among the coloration dictionary's ordered filters:

- **0.0 (bottom)** → darkest filter (strong low-pass character).
- **1.0 (top)** → brightest filter (full high-frequency content).
- The default curve starts at 0.85 and falls to 0.25 — mimicking air absorption (bright-to-dark).
- Flatten it to 1.0 for a consistently bright tail; flatten to 0.0 for a deep, dark reverb throughout.
- Draw a **rise** (dark → bright) for an unusual tonal swell along the tail.

The dictionary contains 6 one-pole lowpass filters ordered dark → bright. Filter routing is per-pulse; the Tone curve value at each pulse's tail position selects the target filter index.

Bottom Strip Dials

Size Tail length, ~0.1 – 10 s. Default: 1.5 s.

Sets the total length of the generated velvet tail. Longer sizes produce larger reverb rooms; combined with a falling Decay curve they give a smooth, room-like decay. Combined with a plateau Decay they produce long, dense sustain.

Design-time parameter — non-automatable. Changing Size re-generates the velvet sequence (an $O(M)$ walk on the GUI thread, not the audio thread). Changes take effect immediately on edit.

Density Pulse density, ~500 – 4000 pulses/s. Default: 1500 /s.

Controls how many velvet pulses are spread per second of tail. Higher density → denser, smoother-sounding tail; lower density → sparser, grainier texture (artefact if too low). The optimal range depends on musical material and the Decay curve.

Design-time parameter — non-automatable. Changing Density re-generates the sequence.

Note: CPU cost scales as **Size × Density** (the sparse-convolution is $O(M)$ per sample). Long tails at high density are expensive; the GUI will still respond but DSP load rises.

Width Maximum right-channel jitter, 0 – 30 ms. Default: 8 ms.

The absolute ceiling for the per-pulse L/R decorrelation offset. The **Width curve** scales this dial value per tail position. At 0 ms the tail is always mono, regardless of the Width curve; at 30 ms the maximum spread reaches 30 ms of physical delay between channels.

Design-time parameter — non-automatable. Changing Width re-generates the sequence.

Pre-Delay Wet-path pre-delay, 0 – 200 ms. Default: 0 ms.

Delays the wet (reverberant) signal relative to the dry signal. The dry signal stays time-aligned; only the reverb tail is pushed back. Use this to simulate room distance — longer pre-delays suggest larger spaces.

Automatable, smoothed.

Mix Dry/wet blend, 0 – 100 %. Default: 30 %.

At 0 % the output is fully dry. At 100 % the output is fully wet (just the tail, no direct signal). Typical room reverb usage: 20–40 %.

Automatable, smoothed.

Input Input trim, –24 – +24 dB. Default: 0 dB.

Scales the signal fed into the reverb engine before velvet convolution. Affects both L and R channels equally. Does not affect the dry path.

Automatable, smoothed.

Output Output trim, –24 – +24 dB. Default: 0 dB.

Scales the final mixed output (after dry/wet blend). Use this to compensate for level changes from Mix or to integrate Nap into a gain-staged mix chain.

Automatable, smoothed.

Seed Integer seed, 1 – 9999. Default: 1.

Selects the random velvet pattern. The generated pulse locations and signs are fully deterministic from the seed and all other design-time parameters. Changing Seed re-rolls the pattern while preserving the curve shapes — the envelope, stereo image, and tonal trajectory all stay the same, but the specific pulse positions shift.

Use this when the current pattern has an audible irregularity or comb artefact. A few seed values usually yield a smooth, natural-sounding tail for any given set of curves.

Design-time parameter — non-automatable. The velvet sequence regenerates automatically whenever the seed (or any other design-time parameter or curve) changes, and after a sample-rate change — there is no manual regenerate step.

Engine Convolution engine selector: **Zero Latency** (default) or **Efficient**.

- **Zero Latency** — time-domain $O(M)$ sparse convolution (the EDVN tapped-delay-line). Reports zero algorithmic latency to the host; the dry and wet paths stay exactly phase-aligned. CPU cost scales linearly with $\text{Size} \times \text{Density}$ (pulse count M).
- **Efficient** — uniformly-partitioned FFT convolution (UPOLS, 512-sample partitions). Nap analytically bakes the wet-path impulse response on the GUI thread and installs it into a `PartitionedConvolver`; the audio thread performs only FFT multiply-accumulate operations rather than M scattered ring-reads per sample. Lower CPU cost at large/dense settings; reports ~ 512 samples (~ 10.7 ms at 48 kHz, proportionally fewer at higher sample rates) of algorithmic latency to the host.

Both engines implement the same EDVN coloration physics — they are mathematically equivalent: $\text{Efficient}[n + 512] \approx \text{ZeroLatency}[n]$ to floating-point precision. Use **Zero Latency** when host latency compensation is unavailable or when absolute minimal delay is required (e.g., live monitoring). Use **Efficient** when recording with PDC and `Size` or `Density` is pushed large enough that the time-domain cost is noticeable.

Switching engines resets the convolver state cleanly and reports the updated latency to the host. Switching to **Efficient** immediately triggers an IR bake if the current sequence is up to date; if a curve drag is in progress, the bake is deferred to drag-release.

Design-time parameter — non-automatable.

How It Works

EDVN Velvet Noise

Velvet noise is a sparse random impulse sequence with exactly one non-zero pulse per fixed grid cell, with a random sign and location within that cell. A velvet-noise reverberator convolves the input with this sparse sequence, producing a reverberation tail whose density is controlled by the grid spacing.

Extended Dark Velvet Noise (EDVN) extends this with a bank of parallel coloration filters (an all-pole “dictionary”) whose job is to give the tail a realistic spectral character — the “dark” quality of real room tails where high frequencies decay faster than low. Each pulse is routed to one dictionary filter; the routing follows the drawn Tone curve, so tonal evolution is under direct user control.

Decay and Coloration Decoupling

A key insight from the EDVN paper (Fagerström et al., JAES 72(6), 2024) is that decay shape and spectral colouring are **mathematically independent**. The energy envelope $g(m)$ applied to each pulse controls level; the filter routing controls colour; they don’t interact. This is why Nap can offer three independently drawn curves — changing Decay doesn’t change Tone, and vice versa.

Width via Pulse-Location Jitter

From the BDVN paper (DAFx 2024), interaural coherence equals the DTFT of the pulse-location jitter probability distribution. A per-pulse right-channel offset $\Delta(m)$ drawn from a distribution of half-width $J(m)$ gives a frequency-dependent coherence curve where wider J produces lower coherence (less correlation between L and R). Making $J(m)$ time-varying — the Width curve — animates the stereo image across the tail at zero added compute cost.

Sparse Convolution (audio thread)

For each output sample:

1. The input sample is written into a pre-allocated ring buffer.
2. Each active velvet pulse tap reads from the ring buffer at its stored offset, scales by its coefficient, and accumulates into one of Q coloration filter inputs.
3. The Q parallel one-pole IIR filters run on their accumulators.
4. Filter outputs are summed $\rightarrow$ a gentle post-filter LP $\rightarrow$ a DC blocker.
5. The result passes through the pre-delay line and is blended with the dry signal.

The left and right channels use the same `coeff[]` and `filter_idx[]` arrays but separate pulse location arrays (`location[]` for L, `location_r[]` for R), so the stereo processing is correct with a single generation pass.

Cost is $O(M)$ per sample where $M = \text{Size} \times \text{Density}$. A 2 s tail at 1500 pulses/s = 3000 taps per sample — well within real-time budget on modern hardware.

Efficient Engine (FFT convolution, optional)

Because the entire wet path is linear and time-invariant, it equals convolution with a single baked impulse response. The **Efficient** engine exploits this:

1. On the GUI thread, `IrBaker` analytically constructs the IR by simulating each coloration filter's response to the velvet pulse field (identical to what `ReverbChannel` computes sample-by-sample), then applies the post-LP and DC blocker in one forward pass over the IR array.
2. The IR is divided into 512-sample partitions, each zero-padded to 1024 points and forward-FFT'd (UPOLS pre-partitioning). The resulting spectra are published to the audio thread via `IrHandoff`.
3. On the audio thread, `PartitionedConvolver` maintains a frequency-domain delay line (FDL) of past input-block spectra. For each new 512-sample input block it runs one forward FFT, accumulates the frequency-domain product against all IR partitions (one multiply-accumulate per partition per bin), and recovers the output via one inverse FFT. The audio thread never transforms the IR itself.

The result is that at large M the $O(M)$ per-sample gather is replaced by $O(k \cdot N \log N)$ per 512-sample block ($k = \text{partition count} \approx \text{IR length} / 512$). At a 10 s $\times$ 4000 /s maximum setting the break-even is well past typical usage, so **Efficient consistently reduces CPU** at the cost of a one-partition (~ 10.7 ms at 48 kHz) algorithmic latency.

Sequence Generation (GUI thread)

The velvet sequence is generated on the GUI thread whenever any design-time parameter or drawn curve changes, then published to the audio thread via a lock-free handoff (`SequenceHandoff`, mirroring `miff's KernelHandoff`). The audio thread calls `try_lock` once per block; if a newer sequence exists, it copies it into its local buffer (pre-allocated to `MAX_PULSES` capacity, no heap allocation) and proceeds.

When the **Efficient** engine is active, any sequence change also triggers an `IrBaker::bake` pass (GUI thread) followed by an `IrHandoff` publish; the audio thread picks up the new spectra on the next block. Continuous MSEG node drags defer the bake to drag-release so the GUI stays responsive during editing; the sequence and the tail visualization still update live on every frame.

Interaction

Drawing Curves

All three panes use the shared MSEG editor from `tiny-skia-widgets`. The curve runs left-to-right over tail position ($0 \rightarrow 100\%$).

- **Double-click empty space** to add a node; **double-click a node** to delete it.
- **Drag a node** to move it. **Left-click and drag on empty space** draws a marquee that selects every node inside it; dragging any selected node then moves the whole group.
- **Drag a segment** (between two nodes) to bend it — this adjusts the segment's tension.
- **Alt-drag** for stepped (staircase) draw mode.
- **Right-click a selected node** opens a transform menu (**Compress** / **Expand** $\times$ **values** / **times**); **right-click anywhere else** toggles that segment's stepped flag.

Any edit triggers immediate re-generation of the velvet sequence. The pulse-field visualization in the Decay pane updates on the next paint frame.

Dial Drag and Text Entry

All bottom-strip dials use the shared `param_dial` widget:

- **Drag vertically** to adjust (up = increase).
- **Shift+drag** for fine control (10 $\times$ slower).
- **Double-click** to reset to default.
- **Right-click** to type an exact value. Press **Enter** to commit, **Escape** to cancel; clicking outside the edit field auto-commits.

Seed

The **Seed** dial selects the integer seed for the velvet pattern. Dragging it re-rolls the pulse locations while keeping the curve shapes. The sequence regenerates automatically on any change to the seed, the design-time parameters, or the curves — and after a sample-rate change — so there is no separate regenerate step.

Resize

Nap's window is freely resizable — drag the plugin window's edge or corner in your host. (Nap has no separate zoom buttons or keyboard shortcuts.) The GUI is CPU-rendered and re-rasterises crisply at any size; the scale tracks the window width (`physical_width / 880`) and clamps to $0.5\times$ – $4.0\times$ (50 % – 400 %). The chosen size is persisted across host restarts via `EditorState`.

Technical Notes

- **Engine-dependent latency** — Zero Latency mode (default): no algorithmic latency reported to the host. Efficient mode: ~ 512 samples (~ 10.7 ms at 48 kHz, fewer at higher sample rates) reported so DAWs with PDC compensate correctly. Pre-Delay is a musical wet-path control in both modes and is not included in the reported latency.

- **No audio-thread allocations** — ring buffers, filter state, and sequence buffers are pre-allocated to maximum capacity; the handoff uses `try_lock` + a bounded copy, never `Vec::new()` or `clone()` in `process()`.
- **No unsafe code** — beyond the standard baseview/raw-window-handle glue inherited from the editor scaffold (where the FFI requires it).
- **CPU rendering** — tiny-skia (software rasterizer) + fontdue (glyph cache) + softbuffer (pixel buffer). No OpenGL context, no GPU driver overhead per instance.
- **Tail reporting** — `process()` returns `ProcessStatus::Tail` while the velvet tail rings out after the input goes silent, so the host knows not to cut processing prematurely.
- **Feature tier** — Nap's O(M) convolver scales with $\text{Size} \times \text{Density}$ and is intentionally not held to the 100+-instance memory target. The DSP is cheap for typical values, but it is a single-instance character reverb, not a metering utility.

Formats

- CLAP
- VST3
- Standalone (JACK, or CPAL — ALSA on Linux, CoreAudio on macOS, WASAPI on Windows — with a dummy/offline fallback)

License

GPL-3.0-or-later