

Satch

User Manual

Michael Dungan

satch Manual




Figure 1: satch interface

What is satch?

satch is a detail-preserving spectral saturator. It uses FFT-based spectral analysis to preserve quiet frequency components (detail) through the clipping process, producing textured flat-top clipping instead of the smooth, featureless flat tops of standard time-domain saturation.

Inspired by Newfangled Audio's Saturate.

Installation

Build from source (requires stable Rust):

```
cargo nih-plugin bundle satch --release
```

The bundler outputs to target/bundled/. Copy either the .vst3 or .clap file to your plugin directory:

- **Linux:** ~/.vst3/ or ~/.clap/
- **macOS:** ~/Library/Audio/Plug-Ins/VST3/ or ~/Library/Audio/Plug-Ins/CLAP/
- **Windows:** C:\Program Files\Common Files\VST3\ or C:\Program Files\Common Files\CLAP\

Quick Start

1. Insert satch on a track
2. Raise **Gain** to push the signal into clipping
3. Turn up **Detail** to preserve quiet harmonics through the clipped sections
4. Adjust **Knee** to taste (0% = hard clip, 100% = soft tanh)

Controls

Row 1: Level

Gain Input gain boost. Range: 0 to +24 dB. Default: 0 dB.

Pushes the signal toward the clip ceiling. Higher gain = more of the signal exceeds the threshold and gets clipped. At 0 dB, the signal passes through at its original level (clipping only occurs if it already exceeds the threshold).

Threshold Clip ceiling. Range: -24 to 0 dB. Default: 0 dB.

Sets where clipping begins. At 0 dB, the ceiling is at full scale (± 1.0). At -6 dB, the ceiling drops to ± 0.5 – any signal above that level is clipped. Below the threshold, the signal passes through unchanged.

Threshold is independent of Gain. Gain controls how hard the signal hits the ceiling; Threshold controls where the ceiling is.

Row 2: Character

Detail Spectral detail preservation. Range: 0 to 100%. Default: 0%.

Controls how much quiet spectral content survives through the clipping process.

- **0%:** Standard time-domain saturation. Flat tops are smooth and featureless – all frequency content in the clipped region is destroyed.
- **100%:** Full spectral detail preservation. Quiet harmonics ride through the clipped sections as visible texture on the flat tops. The peaks stay at the clip level; the texture hangs downward.

Detail only affects clipped portions of the signal. Unclipped material is identical at 0% and 100%.

The spectral analysis uses a 2048-point FFT with 75% overlap (512-sample hop), introducing 2048 samples of latency (compensated automatically by the host).

Knee Clipping shape. Range: 0 to 100%. Default: 100%.

Controls the transition from linear to clipped:

- **0% (hard):** Brick-wall clipping. The signal passes through linearly below the threshold and is hard-clamped at the threshold. Sharp transition, maximum edge.
- **100% (soft):** Tanh saturation curve. Gradual transition into clipping – signals approaching the threshold are gently compressed before reaching the ceiling. Warmer, rounder character.
- **In between:** Linear crossfade between hard and soft.

Bottom Right

Mix Dry/wet blend. Range: 0 to 100%. Default: 100%.

At 100%, the output is fully processed. At 0%, the output is the original dry signal (delayed to match latency). Use intermediate values for parallel saturation.

How It Works

satch combines two clipping paths:

1. **Time-domain path:** Applies $\text{threshold} * \tanh(\text{gain} * x / \text{threshold})$ (crossfaded with a hard clip via the knee) to produce flat-top clipping at the threshold level.
2. **Spectral path:** Uses an STFT (Short-Time Fourier Transform) to split each frame's bins into a *loud* path (bins within 6 dB of the spectral peak) and a *quiet* path (bins more than 20 dB below it, smoothly crossfaded between). The two paths are reconstructed by **separate inverse STFTs**; the nonlinear clipping is applied in the **time domain after reconstruction** (not per-bin in the frequency domain), which keeps the overlap-add normalization exact. The result has flat-top clipping character but with the quiet detail preserved.

The **Detail** knob blends between the two paths, but only in clipped regions. A clip-aware mask ($\tanh^2(\text{gain} * x / \text{threshold})$) detects where clipping is occurring and applies the spectral difference only there. Unclipped material always uses the time-domain path.

The output is safety-clamped at $\pm \text{threshold} \times 1.5$ – detail texture can ride up to 50% above the threshold clip level rather than being limited exactly at it.

Scaling

The window is freely resizable – drag the plugin window's edge or corner in your host. (There are no in-plugin zoom buttons or keyboard shortcuts.) The scale tracks the window width and clamps to 0.5x-4.0x (50% - 400%); the chosen size is persisted across host restarts.

Interaction

- **Drag vertically** on any dial to adjust (up = increase)
- **Shift+drag** for fine control (10x slower)
- **Double-click** any dial to reset to default
- **Right-click** any dial to type a value; **Enter** commits, **Escape** cancels, clicking outside auto-commits

Technical Notes

- **No audio-thread allocations** – the process() callback never allocates heap memory in release builds
- **CPU rendering** – uses tiny-skia (software rasterizer) + fontdue (glyph cache) + soft-buffer (pixel buffer). No OpenGL context, no GPU drivers loaded

- **Time-domain nonlinearity** - the tanh clipping is applied to the reconstructed signal in the time domain (after the two inverse STFTs), *not* per-bin in the frequency domain; this keeps the overlap-add COLA normalization exact. Quiet bins (detail) pass through while the loud fundamental is clipped
- **Clip-aware detail blend** - $\tanh^2$ mask ensures spectral detail is only applied where the signal is actually being clipped, leaving unclipped material identical to the time-domain path
- **Peak-relative spectral split** - bins within 6 dB of the spectral peak are classified as “loud”; bins more than 20 dB below are “quiet”; smooth crossfade in between

Benchmarks (Bitwig, 48 kHz / 1024 samples, GUI closed):

Instances	CPU	RSS	Per Instance
100	13.7%	82 MB	~0.82 MB, 0.14% CPU

Formats

- CLAP
- VST3
- Standalone (JACK, or CPAL - ALSA on Linux, CoreAudio on macOS, WASAPI on Windows - with a dummy/offline fallback)

License

GPL-3.0-or-later